




Implementation and Parallelization of Redfield Equations

CS 205 Final Project
Florian Häse
Sukin Sim
Teresa Tamayo

Background: Exciton Transfer

- Important for improving efficiency of solar cells
- Current methods limited by:
 - small number of excitonic sites (<30) \rightarrow Scaling: $O(N^6)$, N = number of sites
 - short time scales
- Redfield Equation

David Glowacki



Redfield Equation

Bottleneck: matrix-matrix multiplications


$$\frac{d\rho(t)}{dt} = \underbrace{-\frac{i}{\hbar}[H, \rho(t)]}_{\text{System contribution}} + \underbrace{\sum_{k,m,n} \gamma(\omega_{mn}) \left(V_k(\omega_{mn}) \rho(t) V_k^\dagger(\omega_{mn}) - \frac{1}{2} V_k^\dagger(\omega_{mn}) V_k(\omega_{mn}) \rho(t) - \frac{1}{2} \rho(t) V_k^\dagger(\omega_{mn}) V_k(\omega_{mn}) \right)}_{\text{Contribution from coupling between system and environment}},$$

System contribution

Contribution from coupling between system and environment

Dimension: N sites + loss state + target state = N+2 total states

- Dynamics of excited states described by time evolution of density matrix
- Current methods do not scale well beyond systems with > 30 excitonic states.
- Can we do better using different parallelization models? (Long term goal: run a full-scale simulation of large systems)

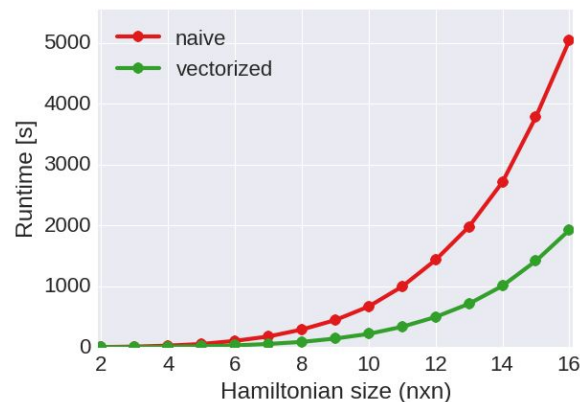



Approaching the problem: Python implementation

1. Identified bottleneck (Lindblad term)


2. Preprocessing the problem ...

- Naive:
 - H , V , ρ complex-valued
- Vectorized:
 - H , V real-valued (can ignore complex phase)
 - H diagonalized to further reduce computational cost
 - V and χ independent of current time step so precomputed and stored in memory for later usage




Increasing Efficiency and Accuracy

- Efficiency: translate code to C to further improve runtimes
- Accuracy: Runge-Kutta 4 (vs. Euler)



Euler integration (Python implementation)


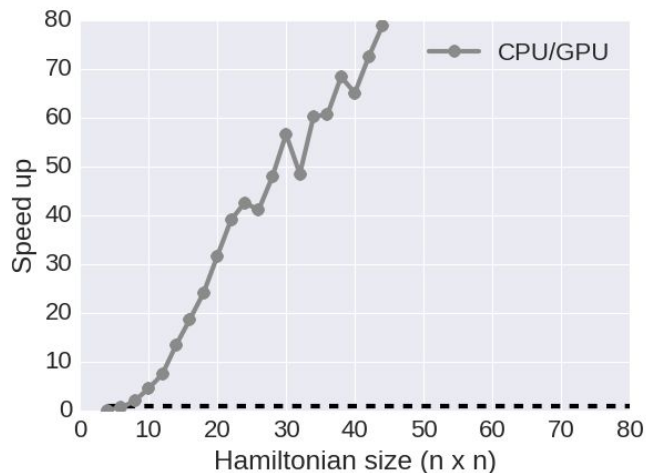



4th order Runge Kutta integration (C implementation)

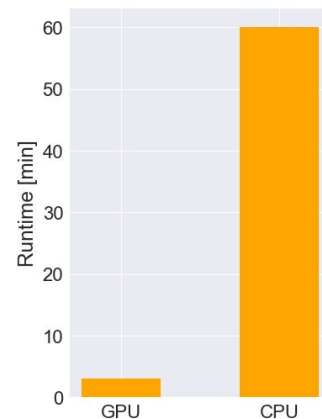
Computing Redfield in SIMT Model: OpenACC on Tesla K80

Features:

- Blocked matrix matrix multiplications
- Reducing the algorithm to the bare essentials



16 site population dynamics:




Parallel Implementation: MPI + OpenAcc

$$\frac{d\rho(t)}{dt} = -\frac{i}{\hbar}[H, \rho(t)] + \sum_{k,m,n} \gamma(\omega_{mn}) \left(V_k(\omega_{mn}) \rho(t) V_k^\dagger(\omega_{mn}) - \frac{1}{2} V_k^\dagger(\omega_{mn}) V_k(\omega_{mn}) \rho(t) - \frac{1}{2} \rho(t) V_k^\dagger(\omega_{mn}) V_k(\omega_{mn}) \right),$$

Contribution from coupling between system and environment

- Current status:
 - Memory Overheads
 - Corruptions in memory



Stronger Scaling via OpenMP

- Bottleneck: matrix-matrix multiplication operations in Lindblad term
 - $15 \cdot N^3$ operations but Runge-Kutta implementation so 4 of these sets per time step
- Approaches: blocking + multithreading
- Better speedups with more threads
- Better scaling as system size increases

